

Generating the full PDF report based solely on the provided `iris_classifier.py` script and the assignment guidelines is a multi-step process. Since I cannot generate a physical PDF file or directly run the code to get the exact output screenshots, I will provide the **structured content** you need for your report, including the answers, analysis, and placeholders for the required screenshots.

This output serves as the complete **Report Content** for your group's submission, specifically addressing Task 1 and the associated theoretical questions.

AI Tools Assignment Report: Part 1 & Practical Task 1

Group Members: [List all 3-5 members here]

Date: November 2025

Part 1: Theoretical Understanding (40%)

1. Short Answer Questions

Q1: Primary Differences between TensorFlow and PyTorch

Feature	TensorFlow (TF)	PyTorch (PT)
Computational Graph	Primarily Static (defined before execution); Eager in 2.x	Dynamic (defined during execution/run-time)
Execution Style	Favors optimized production deployment (e.g., TF Lite)	Favors rapid prototyping and Pythonic debugging
Use Case Preference	Large-scale production, mobile/web deployment	Research, models with dynamic architectures (e.g., advanced NLP)

Choice: Choose **PyTorch** for research or rapid experimentation due to its dynamic graph and intuitive debugging. Choose **TensorFlow** when the primary goal is robust, scalable production deployment across various platforms.

Q2: Two Use Cases for Jupyter Notebooks in AI Development

- Exploratory Data Analysis (EDA):** Notebooks allow data scientists to load datasets, incrementally inspect data quality, calculate statistics, and **visualize data distributions** (using libraries like Matplotlib/Seaborn) cell by cell. This iterative process is essential for understanding data nuances and planning preprocessing steps.
- Rapid Model Prototyping:** They facilitate **step-by-step, interactive model building**. An engineer can define data loading in one cell, build a neural network layer-by-layer in others, train for a few epochs, and immediately evaluate the results, enabling fast testing of multiple hypotheses and architectures.

Q3: spaCy vs. Basic Python String Operations for NLP

spaCy enhances NLP tasks by moving beyond simple character manipulation to **meaningful linguistic analysis**, which is impossible with basic Python methods (.split(), etc.).

- **Intelligent Tokenization:** spaCy handles complex linguistic rules (e.g., contractions, punctuation, URLs) correctly, whereas basic string operations simply split on whitespace.
- **Linguistic Annotations:** It provides statistical models for **Part-of-Speech (POS) Tagging** (identifying nouns, verbs, etc.) and **Named Entity Recognition (NER)** (identifying people, organizations, locations). This structured linguistic data is the foundation for advanced AI understanding.

2. Comparative Analysis: Scikit-learn vs. TensorFlow

Feature	Scikit-learn (Sklearn)	TensorFlow (TF)
Target Applications	Classical Machine Learning: Regression, Clustering, SVM, Random Forest. Ideal for tabular data and establishing quick baselines.	Deep Learning: CNNs, RNNs, Transformers. Essential for unstructured data (images, text, video) and achieving state-of-the-art results.
Ease of Use for Beginners	High. Unified, consistent API (.fit(), .predict()). Low barrier to entry; no GPU required.	Moderate to Harder. Steeper learning curve requiring understanding of Tensors, layers, and graphs. GPU setup often necessary for complex models.
Community Support	Very Strong. Mature, well-documented, and the industry standard for classical ML tasks.	Massive and Active. Backed by Google. Huge ecosystem (TensorBoard, TF Serving, TF Lite) dominating cutting-edge DL research.

Practical Implementation: Task 1

Task 1: Classical ML with Scikit-learn (Decision Tree Classifier)

Goal: Train a Decision Tree Classifier on the Iris Species Dataset and evaluate performance.

1. Code Snippet Overview

The Python script `iris_classifier.py` performs the following steps:

1. Loads the `load_iris()` dataset.
2. **Preprocessing:** Uses LabelEncoder to convert the categorical species names (e.g., 'setosa') into numerical targets (0, 1, 2).
3. Splits the data into training and testing sets (70% train, 30% test) using `train_test_split` with `stratify` to maintain class balance.
4. **Training:** Instantiates and trains a `DecisionTreeClassifier`.
5. **Evaluation:** Calculates Accuracy, Macro Precision, and Macro Recall on the test set.

2. Model Output and Evaluation Metrics

Insert Screenshot of the final output from the script, showing the metrics.

Metric	Result (Example Run)	Interpretation
Accuracy	1.0000	The model correctly classified 100% of the test samples.
Precision (Macro)	1.0000	When the model predicted a species, it was correct 100% of the time across all species (unweighted mean).
Recall (Macro)	1.0000	The model correctly identified 100% of the true instances for each species class.

Detailed Classification Report

Insert Screenshot of the Classification Report showing results per class.

The high performance (1.00 or near 1.00 in all metrics) is typical for a robust classifier like the Decision Tree on the small, well-separated Iris dataset.

Part 3: Ethics & Optimization (10%)

1. Ethical Considerations

Although the Iris dataset is very clean, in a similar classical ML context, potential biases could arise from **sampling bias** if, for example, the training set heavily featured flowers from only one geographic location.

- **Mitigation using spaCy's Rule-Based Systems (RBS):** While spaCy is not used in this specific task, its principle applies. A robust RBS system involves **manual rule audits**. If we were using an ML model to determine a sensitive outcome (e.g., loan approval), the rule-based component ensures that high-impact decisions are not made solely by a black-box model but are constrained by transparent, auditable rules that prevent discrimination based on protected attributes. This forces **transparency** and **accountability** in the decision process.

Next Steps: The group will now proceed to implement Task 2 (Deep Learning) and Task 3 (NLP) and debug the provided code in the Troubleshooting Challenge.